Try again once you are ready.

Required to pass: 80% or higher

You can retake this quiz up to 3 times every 8 hours.

Back to Week 3

Retake

0 / 1 points

1. For which of the following situations would a 0/1 loss function make the most sense?

- Your estimate of the market price of your house.
- Your prediction for the number of bikes sold this year by a local bike shop.

This should not be selected

If the bike shop sells 100 bikes, and you predict sales of 101 bikes, this is just as good of a guess as a prediction of 50 bikes under 0/1 loss.

This question refers to the following learning objective(s):

- Understand the concept of loss functions and how they relate to Bayesian decision making.

- Your prediction as to whether it will rain tomorrow.
- Your answer choice on a Coursera multiple choice quiz.

1 / 1 points

2. Fill in the blank: Under a **linear loss function**, the summary statistic that minimizes the posterior expected loss is the ___________ of the posterior.

- Mode
- Median

Correct

Correct Answer. The median is the summary statistic that minimizes the posterior expected loss under the linear loss function.

This question refers to the following learning objective(s):

- Understand the concept of loss functions and how they relate to Bayesian decision making.

- Mean

---

1 / 1 points

3. You are employed in the Human Resources department of a company and are asked to predict the number of employees that will quit their jobs in the coming year. Based on a Bayesian analysis of historical data, you determine the posterior predictive distribution of the number of quitters to follow a Poisson($\lambda = 10$) distribution. Given a **quadratic** loss function, what is the prediction that minimizes posterior expected loss?

- a. 9
- b. 10

Correct

Correct Answer. 10 is the posterior mean of Poisson($\lambda = 10$). Since the loss function is quadratic, the mean of the posterior distribution minimizes the posterior expected loss.

This question refers to the following learning objective(s):

- Make optimal decisions given a posterior distribution and a loss function.

- c. 11
- d. Either a or b

✖ 0 / 1 points

4. Suppose that you are trying to decide whether a coin is biased towards heads $(p = 0.75)$ or tails $(p = 0.25)$. If you decide incorrectly, you incur a loss of 10. Flipping another coin incurs a cost of 1. If your current posterior probability of a head-biased coin is 0.6, should you make the decision now or flip another coin and then decide?

- Make the decision now, since the minimum posterior expected loss of making the decision now is 4, while the minimum posterior expected loss of making the decision after seeing another coin flip is between 5 and 6.

This should not be selected
Using Bayes rule and the prior, if you observe tails, you should predict a tails bias and if you observe heads, you should predict a heads bias. Use that decision rule to find the posterior expected loss.

This question refers to the following learning objective(s):

- Decide between hypotheses given a loss function: Make optimal decisions given a posterior distribution and a loss function.

- Make the decision now, since the minimum posterior expected loss of making the decision now is 4, while the minimum posterior expected loss of making the decision after seeing another coin flip is between 4 and 5.
- Flip another coin, since the minimum posterior expected loss of making the decision now is 4, while the minimum posterior expected loss of making the decision after seeing another coin flip is between 3 and 4.
- Flip another coin, since the minimum posterior expected loss of making the decision now is 4, while the minimum posterior expected loss of making the decision after seeing another coin flip is between 2 and 3.

---

1 / 1 points

5. You are testing a hypothesis H_1 against an alternative hypothesis H_2 using Bayes Factors. You calculate $BF[H_1 : H_2]$ to be 0.427. According to guidelines first given by Jeffreys (and presented in the lecture), what conclusion can be drawn from the data?

- The data provides strong evidence against H_1.
- The data provides strong evidence against H_2.
- The data provides little to no evidence against H_1.

Correct

Refer to lecture "Posterior probabilities of hypotheses and Bayes factors" to review interpretation of Bayes Factors. The greater $BF[H_1 : H_2]$, the stronger the evidence against H_2.

This question refers to the following learning objective(s):

- Compare multiple hypotheses using Bayes Factors.

- [ ] The data provides significant evidence against H_1.

✔

1 / 1 points

6. Suppose that you are trying to estimate the true proportion p of male births in the United States. Starting with a uniform prior (Beta(1,1)) on p, you randomly sample 10,000 birth certificates, observing 5029 males and 4971 females. What is the posterior probability that p is less than or equal to 0.5? (Hint: use function(s) in R to answer this question)

- [ ] 0.27
- [x] 0.28

Correct
This question refers to the following learning objective(s):

- Conceptualize Lindley's paradox and how the Bayes Factor depends on prior elicitation
- Make inferences about a proportion using a conjugate Beta prior.

- [ ] 0.29

0.30

7. True or False: A Bayesian hypothesis test for a mean $\mu = 0$ using a concentrated prior on μ will yield nearly identical results to a hypothesis test with a high-variance prior μ?

1 / 1 points

True

False

Correct

When testing hypotheses using Bayes Factors, prior specification is highly important. If the prior puts most of its mass on implausible values, the null model will be disproportionately favored.

This question refers to the following learning objective(s):

- Conceptualize Lindley's paradox and how the Bayes Factor depends on prior elicitation
- Compare multiple hypotheses using Bayes Factors.

8. Suppose you are trying to infer the difference in the proportion of yellow M&Ms between two bags of M&Ms. Each bag is estimated to contain roughly 30 M&Ms. To make inferences about the proportion, you draw 10 M&Ms from each bag **without** replacement and record the number of yellow M&Ms in each bag.

0 / 1 points

Is there a problem with the experimental design? If so, what is it?

There is no problem with the experimental design.

This should not be selected

Because we are sampling without replacement and the population size is small, for a given proportion p of yellow M&Ms, our draws are not independent and cannot be modeled b.

This question refers to the following learning objective(s):

- Identify assumptions relating to a statistical inference.

Yes, the probability of drawing a yellow M&M is not independent within groups.

Yes, the probability of drawing a yellow M&M is not independent between groups.

---

1 / 1 points

9. Suppose that when testing $H_0 : p = 0.5$ versus $H_1 : p \neq 0.5$ using Bayes Factors, we get the posterior probability $P(H_0 \mid \text{data}) = 0.25$. Conditional on H_1, the posterior mean of p is 0.6. Under **quadratic** loss, what is the point estimate for p that minimizes expected posterior loss?

0.5

0.55

0.575

Correct

This question refers to the following learning objective(s):

- Create point estimates and credible

intervals by averaging over multiple hypotheses

- Make optimal decisions given a posterior distribution and a loss function.

0.6

---

1 / 1 points

10. True or False: The use of the reference prior $Beta(1/2, 1/2)$ has little bearing on the posterior distribution of a proportion p, provided that the sample size is sufficiently large.

True

Correct
This question refers to the following learning objective(s):

- Conceptualize Lindley's paradox and how the Bayes Factor depends on prior elicitation.

False